1. Select all the correct answers. 1 / 1 point

- [x] With transformers, the vanishing gradient problem isn't related with length of the sequences because we have access to all word positions at all times.

  Correct
  Correct.

- [x] Transformers are able to take more advantage from parallel computing than other RNN architectures previously covered in the course.

  Correct
  Correct.

- [ ] Transformers are models that use both recurrent units and attention mechanisms.

- [x] Even RNN architectures like GRUs and LSTMs don't work as well as transformers for really long sequences.

  Correct
  Correct.

2. Which of the following are applications of transformers? 1 / 1 point

- ( ) Text summarization.
- ( ) Translation
- ( ) Question Answering
- ( ) Chatbots
- (•) All of the above.

  Correct
  Correct. There are others too.

3. What is one of the biggest techniques that the T5 model brings about? 1 / 1 point

- ( ) It's attention mechanism is far more superior than the one used in other models.
- (•) It makes use of transfer learning and the same model could be used for several applications. This implies that other tasks could be used to learn information that would benefit us on different tasks.
- ( ) T5 model is very cheap to train from scratch.
- ( ) It allows for interpretability.

4. When it comes to translating french to english using dot product attention: 1 / 1 point

- [x] You find the distribution by multiplying the queries by the keys (you might need to scale), take the softmax and then multiply it by the values.

Correct
Correct.

- [ ] A CPU is more than enough to train this type of model.
- [x] The intuition is that each query q_i, picks most similar key k_j. This allows the attention model to focus on the right words at each time step.

Correct
Correct.

- [x] The queries are the english words and the keys and values are the french words.

Correct
Correct.

5. Which of the following corresponds to the causal (self) attention mechanism? 1 / 1 point

- ( ) One sentence (decoder) looks at another one (encoder)
- (•) In one sentence, words look at previous words (used for generation). They can not look ahead.
- ( ) In one sentence, in this attention mechanism, words look at both previous and future words.
- ( ) In causal attention, queries and keys come from different sentences and queries search among words before only

Correct
Correct.

6. Let's explore multi-headed attention in this problem. Select all that apply. 1 / 1 point

- [x] Each head learns a different linear transformation to represent words.

Correct
Correct.

- [x] Those linear transformations are combined and run through a linear layer to give you the final representation of words.

Correct

6. Let's explore multi-headed attention in this problem. Select all that apply. 1 / 1 point

- [x] Each head learns a different linear transformation to represent words.

Correct
Correct.

- [x] Those linear transformations are combined and run through a linear layer to give you the final representation of words.

Correct
Correct.

- [x] Multi-Headed models attend to information from different representations at different positions

Correct
Correct.

- [ ] Multi-Headed attention allows you to capture less information than single headed attention.

7. Which of the following is true about about bi-directional attention? 1 / 1 point

- ( ) It only attends to words before.
- ( ) It used an encoder and decodes it using a decoder.
- (•) It could attend to words before and after the target word.
- ( ) It is less powerful than regular uni-directional attention.

Correct
Correct.

8. Why is there a residual connection around each attention layer followed by a layer normalization step in the in the decoder network? 1 / 1 point

- (•) To speed up the training, and significantly reduce the overall processing time.
- ( ) To help with the interpretability.
- ( ) To break the symmetry in the back-prop.
- ( ) To help with the parallel computing component during the training.

Correct
Correct.

Correct.

☐ Multi-Headed attention allows you to capture less information than single headed attention.

7. Which of the following is true about about bi-directional attention? 1 / 1 point

○ It only attends to words before.

○ It used an encoder and decodes it using a decoder.

◉ It could attend to words before and after the target word.

○ It is less powerful than regular uni-directional attention.

Correct
Correct.

8. Why is there a residual connection around each attention layer followed by a layer normalization step in the in the decoder network? 1 / 1 point

◉ To speed up the training, and significantly reduce the overall processing time.

○ To help with the interpretability.

○ To break the symmetry in the back-prop.

○ To help with the parallel computing component during the training.

Correct
Correct.

9. In the lecture, the way summarization is generated is using: 1 / 1 point

○ Next sentence prediction.

○ Next character generation.

◉ Next word generation.

○ By extracting key sentences from the original article.

Correct
Correct.